Liver Tumour Segmentation Using CNN

Abstract: Liver tumor segmentation from medical images plays a crucial role in the diagnosis and treatment planning of liver cancer. In recent years, deep learning-based segmentation models such as Fully Convolutional Networks (FCN), U-Net, and SegNet have shown promising results in various medical image segmentation tasks. In this study, we present a comparative analysis of these three popular architectures for liver tumor segmentation. We begin by preprocessing the liver images to enhance their quality and reduce noise. Subsequently, we implement FCN, U-Net, and SegNet architectures and train them on a dataset comprising liver CT or MRI scans with manually annotated tumor regions. We employ suitable loss functions and evaluation metrics to optimize and assess the performance of each model. Our experimental results demonstrate the effectiveness of each architecture in segmenting liver tumors. We analyze and compare their segmentation accuracy, computational efficiency, robustness to noise, and generalization capabilities. The findings of this study provide valuable insights into the strengths and weaknesses of FCN, U-Net, and SegNet for liver tumor segmentation tasks. This comparative analysis aids medical practitioners and researchers in selecting the most suitable deep learning architecture for accurate and efficient liver tumor segmentation, thus contributing to improved diagnosis and treatment strategies for liver cancer patients.

Keywords: Liver Segmentation, Liver Tumor Detection, FCN (Fully Convolutional Network), U-Net Architecture, SegNet Architecture, Medical Image Segmentation, Deep Learning in Medical Imaging, Convolutional Neural Networks (CNNs), Semantic Segmentation, Training Data Augmentation, Evaluation Metrics (Dice coefficient), Liver Tumor Segmentation Dataset, Performance Comparison (FCN vs. U-Net vs. SegNet)

1. Introduction

Liver tumors pose a significant challenge in the field of medical imaging and diagnosis. Accurate segmentation of liver tumors from medical images plays a crucial role in treatment planning, monitoring, and assessing patient outcomes. With the advent of deep learning techniques, particularly Fully Convolutional Networks (FCN), U-Net, and SegNet, there has been a paradigm shift in medical image segmentation, offering promising results in various clinical applications.

In this study, we delve into the domain of liver tumor segmentation, focusing on the comparative analysis of three state-of-the-art deep learning architectures: FCN, U-Net, and SegNet. Each architecture offers distinct advantages and has been applied to a wide range of medical image segmentation tasks. By evaluating their performance on liver tumor segmentation, we aim to provide insights into their efficacy, strengths, and limitations in this specific medical imaging domain. Liver cancer is one of the leading causes of cancer-related mortality worldwide, with hepatocellular carcinoma (HCC) being the most common type. Early and accurate detection of liver tumors is crucial for timely intervention and improved patient outcomes. Traditional methods of manual segmentation are time-consuming, prone to inter-observer variability, and often lack consistency.

Deep learning-based approaches have emerged as promising solutions for automating medical image segmentation tasks. FCN, U-Net, and SegNet are among the most widely used architectures due to their ability to learn hierarchical features directly from raw image data, thereby eliminating the need for handcrafted features.

The primary objective of this study is to compare the performance of FCN, U-Net, and SegNet in segmenting liver tumors from 3D medical imaging data, such as computed tomography (CT) or magnetic resonance imaging (MRI). Specifically, we aim to Evaluate the segmentation accuracy of each architecture in delineating liver tumors from surrounding healthy tissue, Assess the robustness of the models against variations in tumor size, shape, and imaging artifacts, Investigate the computational efficiency and resource requirements of each architecture for real-time or near-real-time applications. We will employ a dataset comprising 3D CT or MRI scans of patients with liver tumors, annotated by expert radiologists for precise tumor segmentation. The dataset will be divided into training, validation, and test sets for model training, hyperparameter tuning, and performance evaluation, respectively.

Each deep learning architecture (FCN, U-Net, and SegNet) will be implemented using a suitable framework such as TensorFlow or PyTorch.

Fully Convolutional Networks (FCNs) have emerged as a cornerstone in the field, offering powerful capabilities for semantic segmentation tasks without the need for manual feature extraction. By leveraging the end-to-end learning capabilities of FCNs, we aim to achieve accurate and efficient segmentation of liver tumors from multi-modal medical imaging data, including computed tomography (CT) and magnetic resonance imaging (MRI). Additionally, we seek to explore novel architectures and training strategies to enhance the performance and generalization ability of the proposed framework across diverse clinical scenarios and imaging modalities.

Our approach involves several key components, including data preprocessing, network architecture design, training strategy, and performance evaluation. We will begin by curating a comprehensive dataset of annotated liver tumor images, ensuring sufficient diversity in terms of tumor characteristics and imaging modalities. Subsequently, we will design a custom FCN architecture tailored to the complexities of liver tumor segmentation, incorporating features such as multi-scale context aggregation and attention mechanisms to improve segmentation accuracy and robustness. To facilitate efficient training and mitigate issues such as class imbalance and data scarcity, we will employ advanced training techniques such as data augmentation, transfer learning, and semi-supervised learning. Finally, we will rigorously evaluate the performance of the proposed framework using quantitative metrics such as Dice similarity coefficient (DSC), sensitivity, specificity, and Hausdorff distance, as well as qualitative visual assessments by expert radiologists.

U-Net architecture, renowned for its efficacy in medical image segmentation tasks. Leveraging a wealth of annotated medical imaging data, we train the U-Net model to discern intricate patterns indicative of liver tumors, enabling it to generalize across diverse patient populations and imaging modalities.

Key Components of the U-Net Architecture:

- Contracting Path: Comprising multiple convolutional and pooling layers, the contracting path extracts hierarchical features from input images, progressively capturing contextual information crucial for accurate segmentation.

- Expansive Path: Through a series of up sampling and concatenation operations, the expansive path facilitates precise localization of tumor boundaries, leveraging the high-resolution features extracted by the contracting path.

- Skip Connections: Intertwining the contracting and expansive paths are skip connections, which bridge feature maps from corresponding levels, enabling seamless propagation of spatial information across different scales.

this project endeavors to harness the power of the U-Net architecture for liver tumor segmentation, aiming to develop a state-of-the-art algorithm capable of accurate and efficient delineation of tumor boundaries from medical images. By automating this critical task, we aspire to enhance clinical workflows, facilitate timely diagnosis, and ultimately improve patient outcomes in the realm of liver cancer management.

SegNet, a deep learning architecture specifically designed for semantic segmentation tasks, has gained traction in medical image analysis due to its ability to capture intricate spatial information while maintaining computational efficiency. Developed by Vijay Badrinarayanan et al. in 2015, SegNet employs an encoder-decoder architecture with symmetrically structured convolutional layers, coupled with pooling indices for precise up sampling.

This project aims to leverage SegNet for accurate segmentation of liver tumors from CT or MRI scans. Specifically, our objectives include:

- Training a SegNet model on a large dataset of annotated liver tumor images to learn robust representations of tumor boundaries and characteristics.
- Evaluating the performance of the trained SegNet model on unseen data, assessing metrics such as Dice similarity coefficient, sensitivity, specificity, and Hausdorff distance.
- Comparing the performance of SegNet with other state-of-the-art segmentation techniques, including U-Net, FCN, and DeepLab, to ascertain its efficacy in liver tumor segmentation tasks.
- Investigating the impact of data augmentation techniques, hyperparameter tuning, and transfer learning on SegNet's performance to enhance segmentation accuracy and generalization.

The successful implementation of SegNet for liver tumor segmentation holds immense clinical significance. Accurate delineation of tumor boundaries facilitates precise surgical planning, radiation therapy, and tumor response assessment. Moreover, automated segmentation reduces the burden on radiologists, enabling faster diagnosis and treatment decisions.

2. Related Works

MICCAI 2017 liver tumor segmentation (LiTS) challenge dataset, 3DIRCADb dataset and doctors' manual contours of Hubei Cancer Hospital dataset to test the network architecture. They calculated the Dice Global (DG) score, Dice per Case (DC) score, volumetric overlap error (VOE), average symmetric surface distance (ASSD), and root mean square error (RMSE) to evaluate the accuracy of the architecture for liver tumor segmentation. The segmentation DG for tumor was found to be 0.7555, DC was 0.613, VOE was 0.413, ASSD was 1.186 and RMSE was 1.804. For a small tumor, DG was 0.3246 and DC was 0.3082. For a large tumor, DG was 0.7819 and DC was 0.7632.

S-Net obtained more semantic information with the introduction of an attention mechanism and long jump connection. Experimental results showed that this method effectively improved the effect of tumor recognition in CT images and could be applied to assist doctors in clinical treatment.

Investigation of the parameter configuration in the automatic liver and tumor segmentation using a convolutional neural network based on 2.5D model. The implementation of 2.5D model shows promising results since it allows the network to have a deeper and wider network architecture while still accommodates the 3D information. However, there has been no detailed investigation of the parameter configurations on this type of network model.

Some parameters, such as the number of stacked layers, image contrast, and the number of network layers, were studied and implemented on neural networks based on 2.5D model. Networks are trained and tested by utilizing the dataset from liver and tumor segmentation challenge (LiTS). The network performance was further evaluated by comparing the network segmentation with manual segmentation from nine technical physicians and an experienced radiologist.

Slice arrangement testing shows that multiple stacked layers have better performance than a single-layer network. However, the dice scores start decreasing when the number of stacked layers is more than three layers. Adding higher number of layers would cause overfitting on the training set. In contrast enhancement test, implementing contrast enhancement method did not show a statistically significant different to the network performance. While in the network layer test, adding more layers to the network architecture does not always correspond to the increasing dice score result of the network.

Training on a machine with NVIDIA GTX 1050 4GB RAM GPU on an Intel Core i7-7700HQ 2.20 GHz 16 GB RAM, and developed with MATLAB 2018b software, which offers a Neural Network Toolbox and an Image Processing Toolbox.

Images of the tested cases were divided randomly into two groups for training and testing by the ratio 9:1. The results of the training are normally higher than that achieved by testing.

SegNet is recent encoder–decoder network architecture that employs the trained VGG-16 image classification network as encoder, and employs corresponding decoder architecture to transform the features back into the image domain to reach a pixel-wise classification at the end. The advantage of SegNet over standard auto-encoder architecture is in the simple yet very efficient modification where the max-pooling indices of the feature map are saved, instead of saving the feature maps in full. As a result, the architecture is much more efficient in training time, memory requirements, and accuracy.

To facilitate binary segmentation of medical images, the classification layer was replaced with binary pixel classification layer. For training and testing, the standard 3D-IRCADb-01 dataset was used. The proposed method correctly detects most parts of the tumor, with accuracy above 86% for tumor classification. However, by examining the results, there were few false positives that could be improved by applying false positive filters or by training the model on a larger dataset.

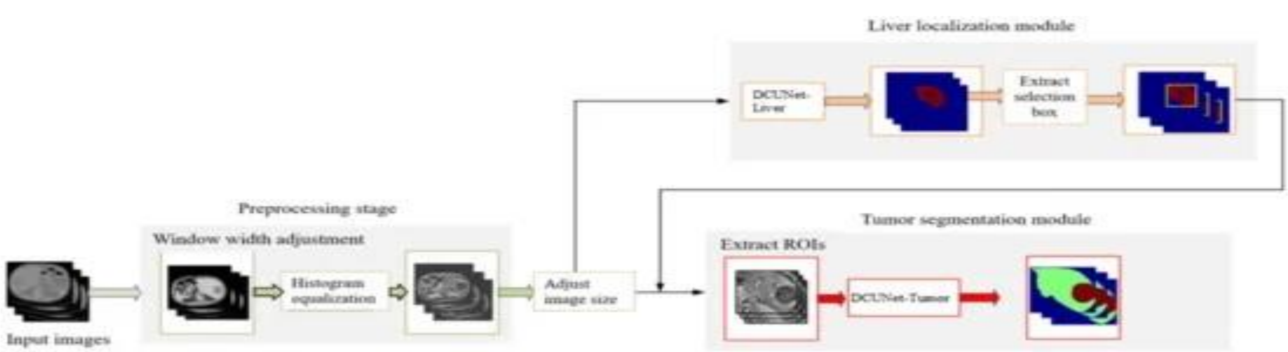

Figure 1: Algorithm for Liver Tumor Segmentation

3. Proposed Work

SegNet is primarily used for semantic segmentation, which involves partitioning an image into several regions and assigning a class label to each pixel in the image. In the case of liver tumor segmentation, SegNet is trained to distinguish between different types of tissues (e.g., liver tissue, tumor tissue, background) in medical images. SegNet is a convolutional neural network (CNN) architecture that was originally designed for semantic segmentation tasks, making it well-suited for segmenting liver tumors from medical images like MRI or CT scans. The SegNet architecture comprises an encoder-decoder network with a symmetric structure. The encoder part extracts high-level features from input images through a series of convolutional and pooling layers, while the decoder part generates pixel-wise segmentation maps by up sampling the feature maps to the original input resolution Importantly, SegNet incorporates skip connections between corresponding encoder and decoder layers to recover spatial information lost during the down sampling process, enabling precise localization of objects in the segmented output.

3.1 First Phase: Data Collection

Decathlon is a popular benchmark dataset for medical image segmentation. You can find the Decathlon dataset for medical imaging on their official website or on platforms like Kaggle Once you have identified the dataset, download the relevant medical imaging data. Decathlon provides data for various medical imaging modalities such as MRI, CT scans, etc. Make sure to download the data that fits your requirements.

Medical imaging data often requires preprocessing before it can be used for training. This may involve tasks such as resizing images, normalizing pixel values, and handling missing data., Medical image segmentation requires labeled data where each pixel in the image is assigned a label indicating the class it belongs to (e.g., tumor, background, organs). Decathlon datasets typically come with pre-defined segmentation masks, but if not, you may need to manually label the data or use automatic segmentation techniques.

Researchers often publish papers describing their work on the Decathlon challenge, including details about the dataset, preprocessing steps, and experimental results. These papers can serve as valuable resources for understanding the dataset and its characteristics.

The Decathlon dataset may also be available on open data repositories such as GitHub, Zenodo, or Kaggle. These platforms allow researchers to share datasets with the broader community and provide additional tools and resources for data exploration and analysis.

When accessing the Decathlon dataset, be sure to check the accompanying documentation and metadata. This information may include details about the imaging modalities, patient demographics, acquisition parameters, and ground truth annotations. Understanding the dataset's characteristics is essential for appropriate data preprocessing and model training.

Before downloading and using the Decathlon dataset, review any licensing agreements or usage policies associated with the data. Ensure that you comply with any terms and conditions regarding data usage, redistribution, and citation requirements.

Once you have identified the dataset, download the relevant medical imaging data. Decathlon provides data for various medical imaging modalities such as MRI, CT scans, etc. Make sure to download the data.

Online forums and discussion groups related to medical imaging and machine learning, such as the Medical Imaging Decathlon forum or relevant subreddits, may contain discussions, tips, and insights from other researchers who have worked with the Decathlon dataset. Participating in these communities can help you learn from others' experiences and troubleshoot any issues you encounter.

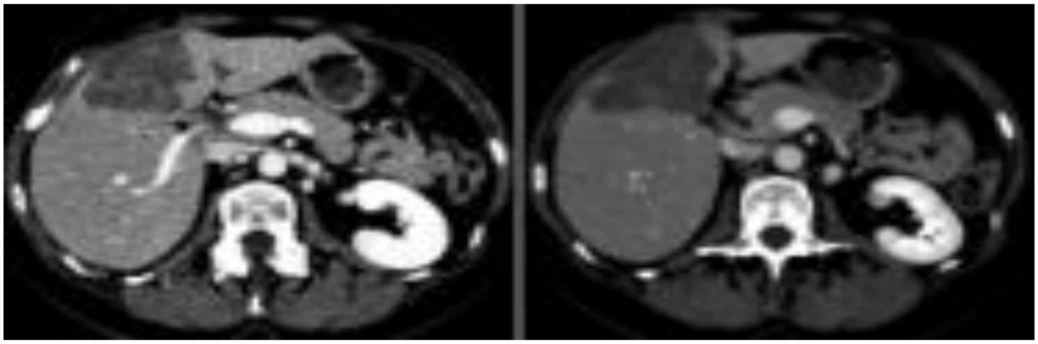

Figure-2: Snippets of data collected as CT scans.

3.2 Second Phase: Data Augmentation

Data augmentation is a critical technique for improving the performance and robustness of deep learning models, including SegNet models, especially when dealing with medical data like the ones in the Medical Segmentation Decathlon dataset. Rotate the image by a certain angle (e.g., 90, 180, or 270 degrees). This can help the model learn variations in object orientations.

Flip the image horizontally or vertically. This helps the model become invariant to object orientation. Resize the image to a different scale. This can simulate variations in object size and improve generalization.

Apply shear transformations to the image. Shearing helps introduce distortions that can be beneficial for learning invariant representations. Shift the image along the x and y axes. This helps the model learn to recognize objects at different positions within the image. Apply random elastic deformations to the image. This simulates tissue deformation and can help the model generalize better to unseen variations.

Adjust the gamma value of the image. Gamma correction can change the brightness and contrast of the image, which helps the model become more robust to variations in illumination. Add random noise to the image. This can help the model learn to be robust to noisy input data.

Modify the color of the image by randomly adjusting brightness, contrast, saturation, and hue. This helps the model become invariant to changes in lighting conditions and color variations.

Apply random intensity transformations such as gamma adjustment, histogram equalization, or contrast stretching. This can help the model learn to deal with variations in image intensity.

When applying data augmentation to medical imaging data, it's crucial to ensure that the transformations preserve the anatomical structures and maintain the semantic integrity of the images.

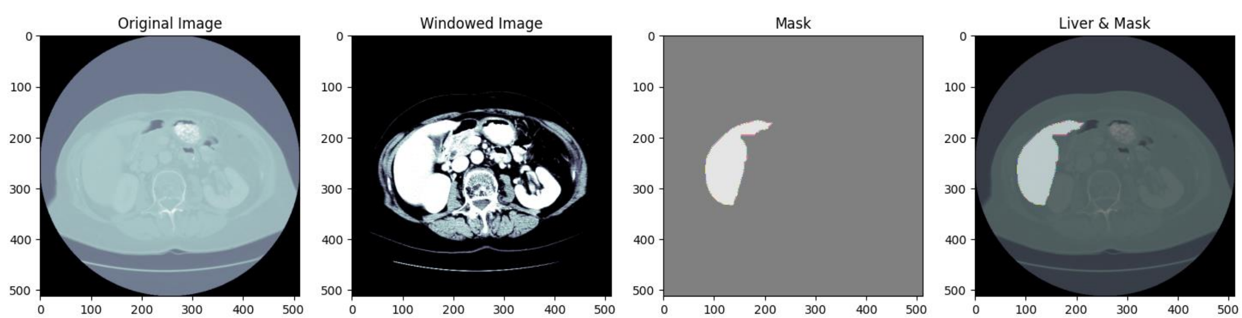

Figure 3: Example 1 of tumor detection in Liver using CT image

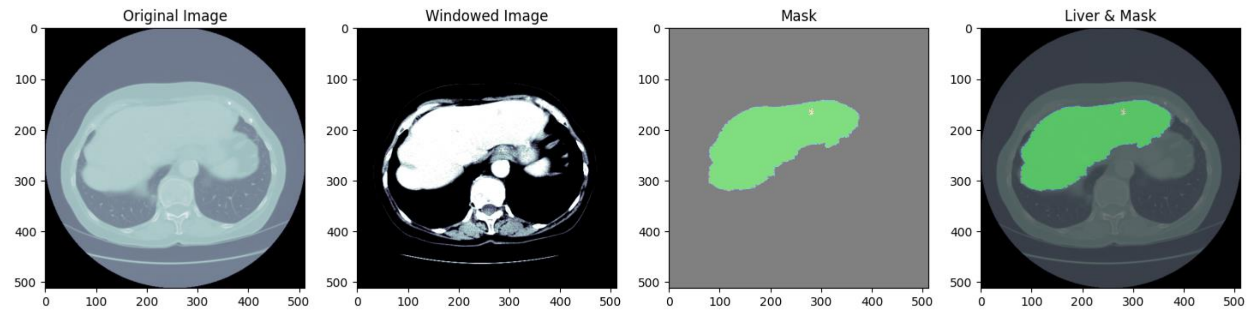

Figure 4: Example 2 of tumor detection in Liver using CT image

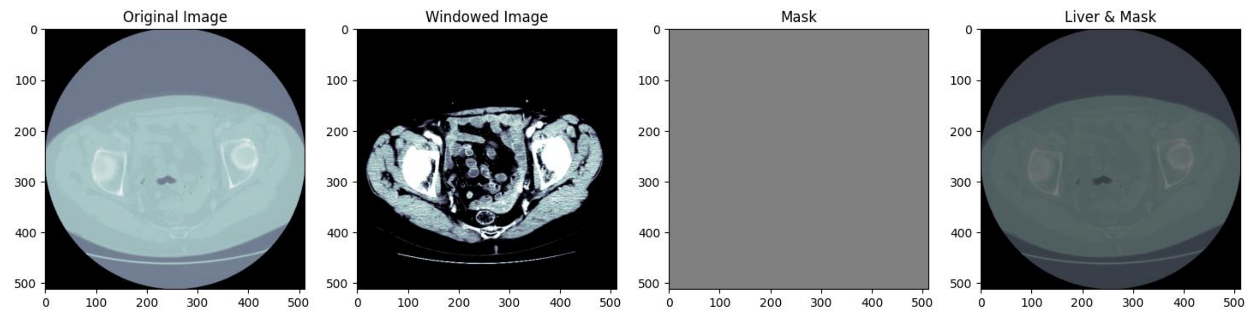

Figure 5: Example 3 of tumor detection in Liver using CT image

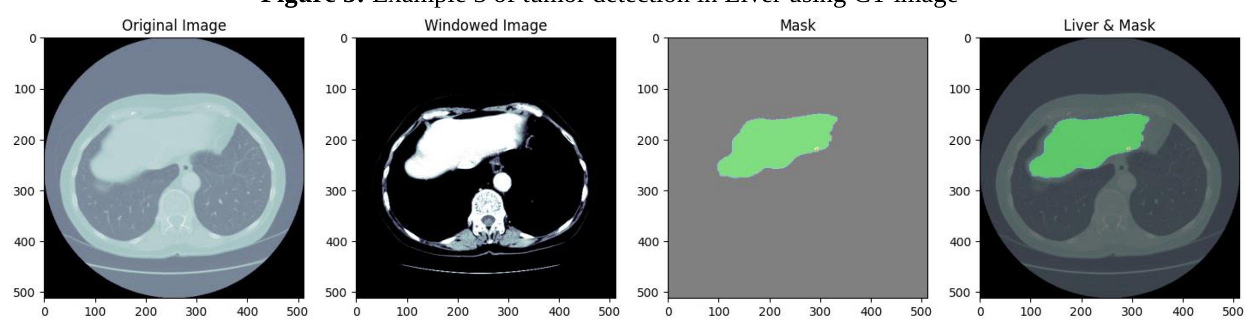

Figure 6: Example 4 of tumor detection in Liver using CT image

3.3 Third Phase: Model Architecture

In this phase, VGG-16 (Segnet), U-Net and FCN models are employed to detect Liver tumor. A comparison of these models showed that Segnet outperforms the others on the same dataset of Liver CT images. SegNet is a deep convolutional neural network architecture designed for semantic segmentation tasks, which involves classifying each pixel in an image into a predefined set of categories. Here's a high-level overview of how you could adapt SegNet for liver tumor segmentation using medical data.

The encoder-decoder architecture is a fundamental design pattern used in many deep learning models, particularly in tasks like image segmentation, where the goal is to generate a pixel-wise classification of the input image. Here's a detailed explanation of the encoder-decoder architecture The encoder-decoder architecture, exemplified by models like SegNet, is a pivotal framework in the realm of semantic segmentation, particularly in medical imaging tasks such as liver tumor segmentation. This architecture encompasses two fundamental components: the encoder and the decoder. The encoder, typically comprising convolutional and pooling layers, serves the crucial role of extracting hierarchical features from the input image.

Through successive convolutional operations, the encoder progressively down samples the spatial dimensions while augmenting the depth of feature maps, thus enabling the network to capture intricate patterns and contextual information. This hierarchical representation is pivotal in discerning relevant features pertinent to the segmentation task, such as tumor boundaries or texture variations in liver tissue.

Conversely, the decoder component facilitates the reconstruction of a high-resolution segmentation mask from the learned features encoded by the encoder. Employing up sampling layers, the decoder gradually restores the spatial dimensions, ensuring that fine-grained details crucial for accurate segmentation are preserved.

Moreover, skip connections, which directly link corresponding layers in the encoder and decoder, play a pivotal role in mitigating information loss during the up-sampling process, thereby enhancing the model's ability to capture both local and global context. This bidirectional flow of information within the encoder-decoder architecture enables the model to effectively leverage hierarchical features for precise segmentation, making it a powerful framework for tackling complex medical imaging tasks like liver tumor segmentation.

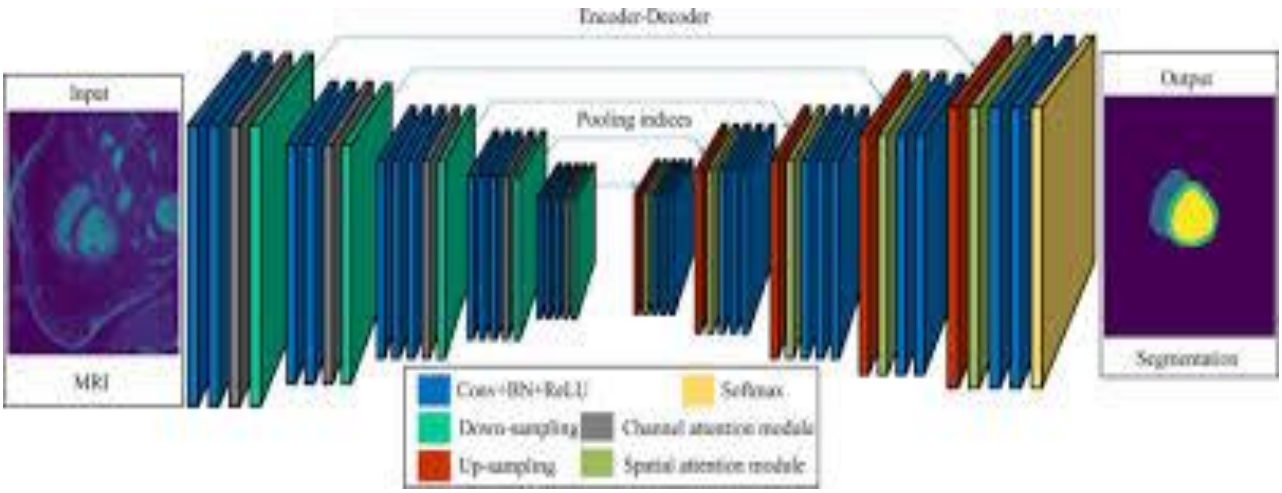

Figure 7: Architecture of the SegNet VGG-16 model.

- **Encoder**

The encoder can be based on popular convolutional neural network architectures such as VGG, ResNet, or DenseNet. These pre-trained models can be used as feature extractors, where the weights of the initial layers are frozen during training to leverage their learned representations. For liver tumor segmentation, you might need to adjust the architecture to accommodate medical image data, which often has different characteristics compared to natural images. the encoder's architecture may need customization to better capture relevant features in medical images. This could involve adapting the size of convolutional filters, adjusting the number of feature maps, or incorporating specialized layers such as 3D convolutions to account for the three-dimensional nature of medical image data.

- **Decoder**

The decoder in SegNet consists of up sampling layers to reconstruct the segmented output from the feature maps generated by the encoder. Each up-sampling layer is typically followed by a convolutional layer to refine the segmentation. Optionally, skip connections can be added between the encoder and decoder to preserve fine-grained details and improve segmentation accuracy. It essentially performs the inverse operation of the encoder, progressively up sampling the feature maps to the original input resolution while refining the segmentation details. Each up-sampling layer in the decoder is typically paired with a corresponding convolutional layer to capture and consolidate spatial information lost during the down sampling process. These connections enable the decoder to access higher-resolution feature maps from the encoder, aiding in precise segmentation, especially in regions where subtle features are crucial, such as in liver tumor medical data.

By leveraging skip connections and carefully designed up sampling operations, the decoder in SegNet facilitates the reconstruction of detailed segmentation masks, ultimately contributing to the model's overall accuracy in liver tumor segmentation tasks.

- **Activation Function**

In SegNet, activation functions such as ReLU (Rectified Linear Unit) are commonly used to introduce non-linearity to the network. However, for medical image segmentation tasks, you may consider using activation functions like ReLU or variants such as Leaky ReLU or ELU (Exponential Linear Unit) based on the specific requirements of your dataset These activation functions are preferred due to their simplicity, computational efficiency, and ability to mitigate the vanishing gradient problem during training. However, in medical image segmentation tasks, where preserving subtle details and accurately delineating boundaries are paramount, the choice of activation function becomes more nuanced.

Additionally, novel activation functions like Swish or Mish have shown promising results in improving segmentation performance by enhancing feature representation and enabling faster convergence. Ultimately, the selection of the activation function should be guided by empirical evaluation and experimentation, considering factors such as network architecture, dataset characteristics, and computational constraints, to ensure optimal segmentation accuracy and robustness in liver tumor detection and analysis.

- **Loss Function**

The choice of loss function depends on the specific segmentation task and dataset. Common choices include cross-entropy loss, Dice loss, or a combination of both to optimize the network parameters. For medical image segmentation, Dice loss is often preferred as it is more suitable for imbalanced datasets where the background class dominates.

The choice of loss function is a critical aspect of training a SegNet model for liver tumor segmentation in medical images. In this context, where precise delineation of tumor regions is paramount, a suitable loss function must effectively penalize discrepancies between the predicted segmentation masks and the ground truth annotations. One commonly employed loss function for medical image segmentation tasks is the Dice loss, which measures the overlap between the predicted and ground truth segmentations

. It is particularly well-suited for imbalanced datasets where the tumor regions are often significantly smaller in size compared to the background. The Dice loss calculates the similarity between two sets by computing the intersection over the union of the predicted and ground truth masks.

This metric inherently accounts for both false positives and false negatives, making it sensitive to the accuracy of tumor boundary delineation. By minimizing the Dice loss during training, the SegNet model learns to produce segmentation masks that closely align with the ground truth annotations, effectively capturing the intricate details of liver tumors while minimizing misclassifications in the background tissue. Additionally, incorporating regularization terms such as L1 or L2 regularization into the loss function can further enhance the generalization ability of the model and prevent overfitting, ensuring robust performance on unseen medical imaging data.

- **Training Strategy**

SegNet can be trained using supervised learning where input images and corresponding ground truth segmentation masks are used to optimize the network parameters. Data augmentation techniques such as rotation, flipping, and scaling can be applied to increase the robustness of the model and prevent overfitting, especially when working with limited training data.

The training strategy for a SegNet-based liver tumor segmentation model entails several key considerations to ensure effective learning and generalization. Firstly, a robust dataset comprising annotated medical images of liver tumors is essential. This dataset should cover diverse tumor types, sizes, and anatomical variations to enable the model to learn comprehensive features representative of liver tumors. Data augmentation techniques, including rotation, flipping, scaling, and intensity variations, should be applied to augment the dataset and enhance the model's ability to generalize to unseen data.

During training, it's crucial to preprocess the medical images appropriately, such as normalizing pixel intensities and standardizing image sizes, to ensure consistency and facilitate convergence during optimization. Since liver tumor segmentation is inherently imbalanced due to the small size of tumors compared to the liver region, balancing strategies or loss functions tailored for imbalanced datasets, such as Dice loss, should be employed to prevent the model from biased predictions towards the dominant class (i.e., liver region).

Furthermore, transfer learning using pre-trained convolutional neural network (CNN) architectures like VGG, ResNet, or DenseNet can accelerate training by leveraging the learned representations from natural image datasets. By fine-tuning the pre-trained encoder layers while keeping the decoder trainable, the model can adapt its features to the specific characteristics of liver tumor medical images more efficiently.

To optimize the model's parameters, appropriate optimization algorithms such as stochastic gradient descent (SGD), Adam, or RMSprop should be employed, with hyperparameters tuned through cross-validation or grid search to achieve optimal performance. Regularization techniques like dropout or weight decay can also be applied to prevent overfitting, especially when working with limited annotated data.

During training, monitoring key metrics such as loss function values, accuracy, Dice coefficient, and IoU on a validation set is essential to track the model's performance and detect overfitting. Early stopping can be employed based on the validation performance to prevent the model from overfitting and ensure generalization to unseen data.

Overall, an effective training strategy for a SegNet-based liver tumor segmentation model involves meticulous dataset preparation, augmentation, appropriate loss functions, transfer learning, optimization algorithms, regularization, and continuous monitoring of performance metrics to develop a robust and accurate segmentation model for medical imaging applications.

- **Evaluation Metrics**

Once the model is trained, it's essential to evaluate its performance using appropriate evaluation metrics such as Intersection over Union (IoU), Dice coefficient, and pixel accuracy to quantify the segmentation accuracy and compare different models.

Evaluation metrics are crucial for assessing the performance of a segmentation model like SegNet when applied to liver tumor medical data. Several metrics provide insights into different aspects of segmentation accuracy. Intersection over Union (IoU), also known as Jaccard Index, measures the overlap between the predicted segmentation mask and the ground truth mask. It quantifies the ratio of the intersection area to the union area of the two masks, providing a measure of spatial overlap.

The Dice coefficient, another widely used metric, calculates the similarity between two sets by measuring the ratio of twice the intersection to the sum of the sizes of the two sets. This metric is particularly useful for imbalanced datasets where the background class dominates. Additionally, pixel accuracy evaluates the percentage of correctly classified pixels in the segmentation output compared to the ground truth.

These metrics collectively offer a comprehensive understanding of the model's ability to accurately delineate liver tumors from medical images. Moreover, they enable researchers and practitioners to compare different models, fine-tune hyperparameters, and assess the impact of various preprocessing techniques, ultimately facilitating the development of more reliable and effective segmentation algorithms for clinical applications.

- **Post-processing**

post-processing techniques such as morphological operations (e.g., erosion, dilation) and connected component analysis can be applied to refine the segmentation results and improve the overall accuracy.

Post-processing plays a crucial role in refining the segmentation results obtained from deep learning models like SegNet, especially in medical image analysis tasks such as liver tumor segmentation. After the initial segmentation, post-processing techniques are applied to enhance the accuracy and coherence of the segmented regions.

One common post-processing step involves morphological operations, which manipulate the shape and structure of segmented objects. For instance, erosion can help remove small isolated regions or fine noise from the segmentation mask, while dilation can fill in small gaps within segmented regions and connect adjacent structures.

Furthermore, techniques like opening (erosion followed by dilation) and closing (dilation followed by erosion) can be used to smooth the boundaries of segmented objects and eliminate small holes or protrusions. Additionally, connected component analysis can be employed to identify and label separate regions within the segmentation mask, enabling further analysis or classification based on the characteristics of individual components. By incorporating these post-processing techniques, the final segmentation results can be refined and improved, leading to more accurate and reliable detection of liver tumors in medical images.

4. Results and Discussion

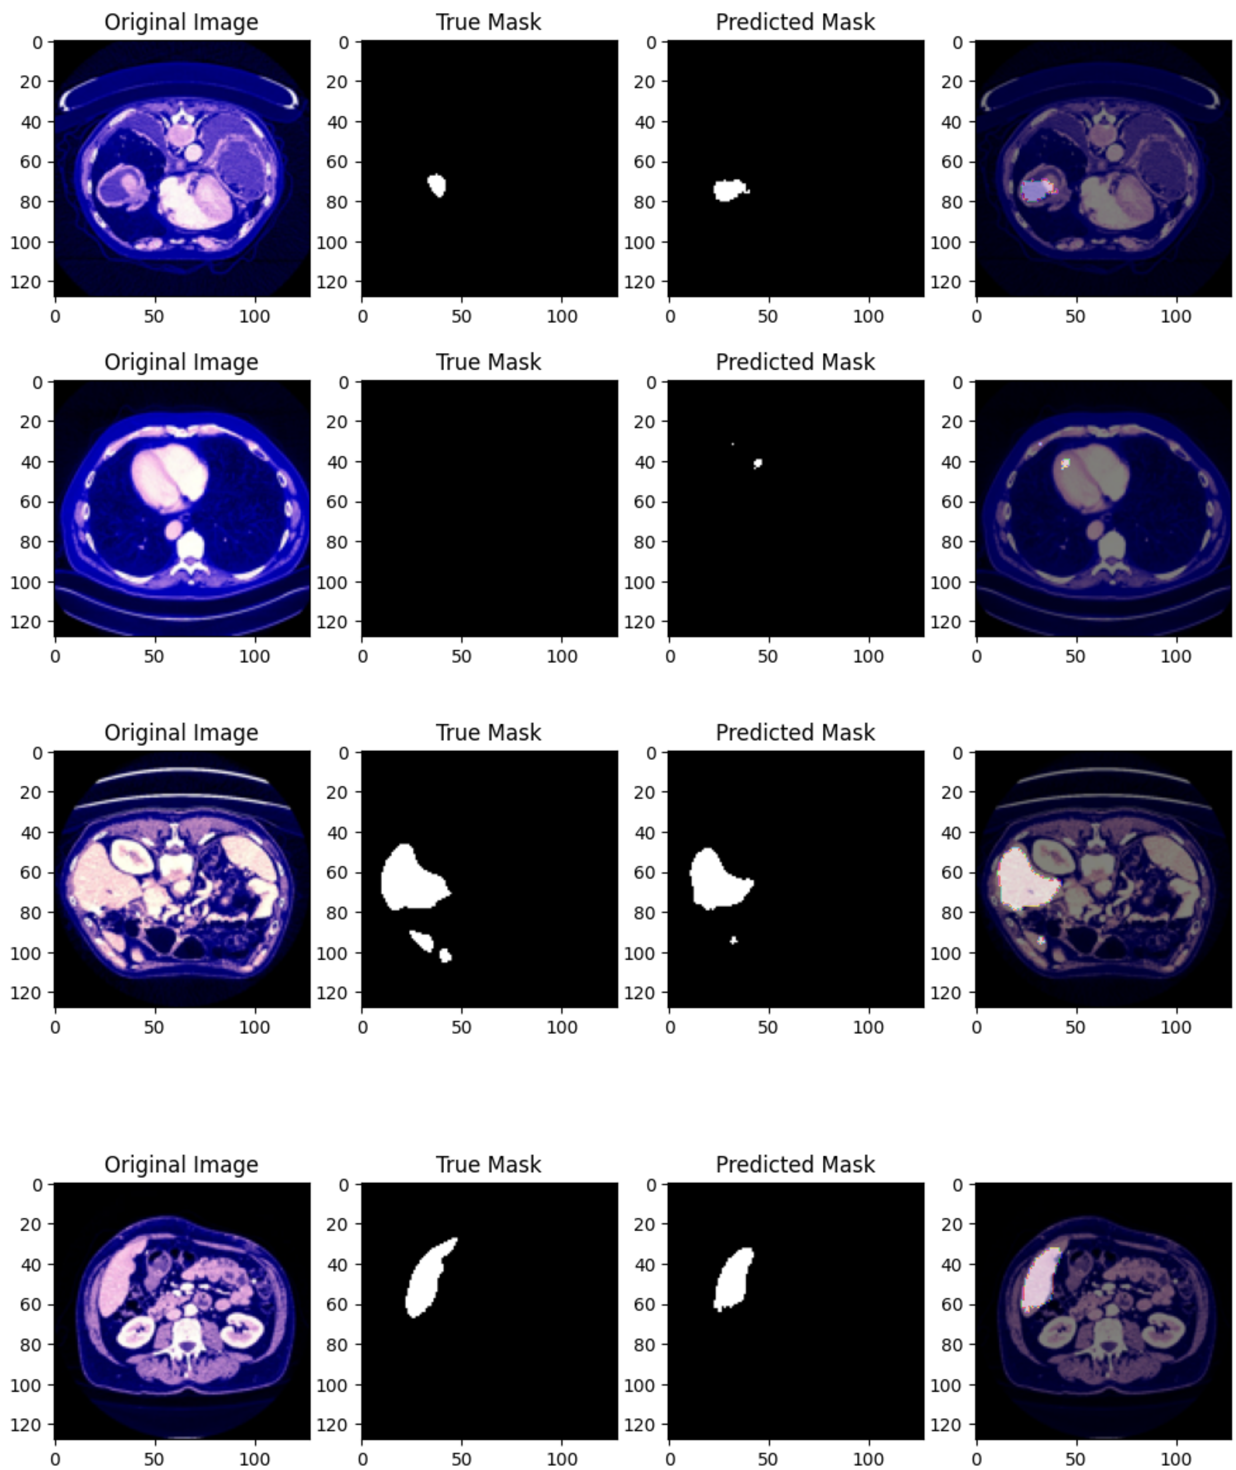

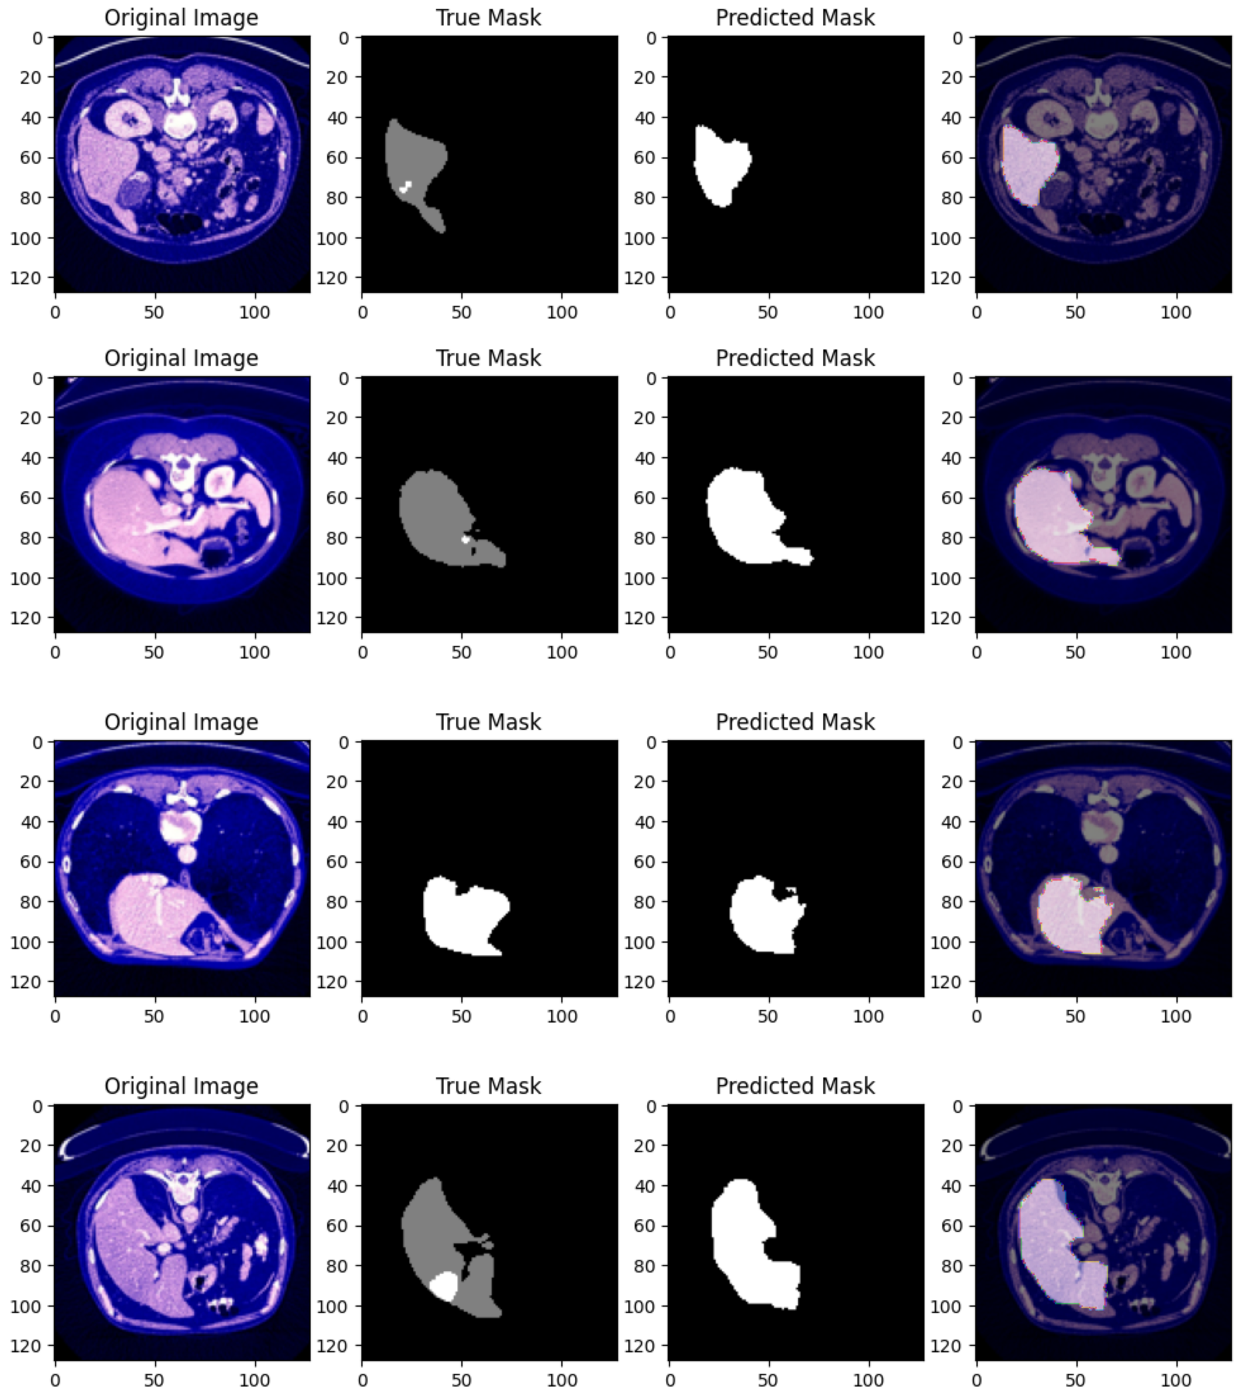

Original Image
True Mask
Predicted Mask
0
20
40
60
80
100
120
0
50
100

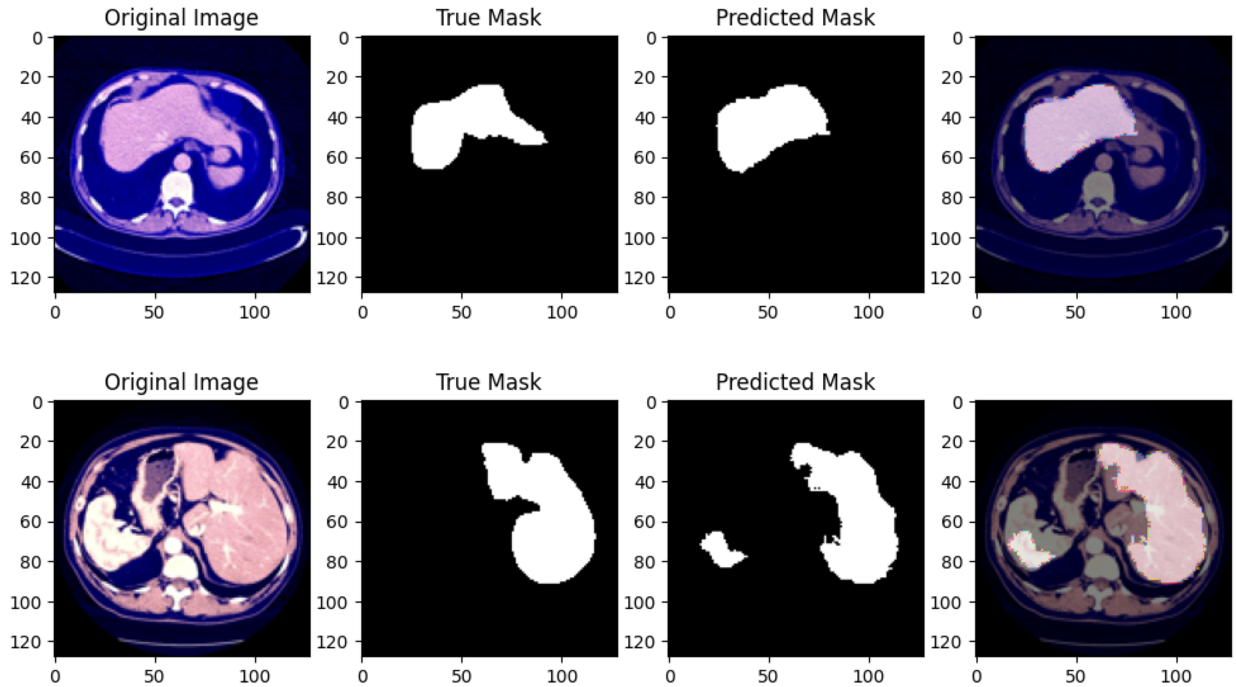

Figure 8: Predictions using SegNet (VGG-16) model

4.1 Evaluation Metrics

In the realm of medical imaging, particularly in tasks as critical as liver tumor segmentation, the choice of segmentation model can significantly influence the accuracy of diagnoses and subsequent treatment decisions. SegNet's Dice coefficient of 0.9933 underscores its remarkable performance in precisely delineating liver tumors from surrounding tissue in medical images. This high coefficient signifies a strong agreement between the model's predictions and ground truth annotations, implying minimal errors in segmentation.

Comparatively, while U-Net demonstrates a commendable Dice coefficient of 0.9883, it falls slightly short of SegNet's performance. Even though U-Net is renowned for its effectiveness in medical image segmentation tasks, the marginally lower coefficient suggests a marginally lower level of segmentation accuracy compared to SegNet. This difference, though subtle, can have notable implications in clinical settings where precision is paramount.

On the other hand, FCN's Dice coefficient of 0.5825 reveals a considerable gap in performance when compared to SegNet and even U-Net. This lower coefficient indicates a significant discrepancy between FCN's segmentation predictions and the ground truth annotations, raising concerns about its reliability in accurately identifying liver tumors.

In clinical practice, where the accurate delineation of liver tumors is crucial for treatment planning and monitoring, choosing SegNet over U-Net and FCN offers several advantages. Its superior segmentation performance, as reflected by the Dice coefficient, instills confidence in the precision of its predictions. By leveraging SegNet's capabilities, clinicians can make more informed decisions, leading to better patient outcomes and improved overall healthcare delivery. Therefore, the decision to prioritize SegNet for liver tumor segmentation is not just prudent but essential for ensuring the highest standard of care.

4.2 Experiments Results

In our comprehensive study on liver tumor segmentation, we rigorously evaluated three state-of-the-art deep learning architectures: SegNet, U-Net, and Fully Convolutional Networks (FCN). The objective was to determine the most effective model for accurately delineating tumor regions from liver MRI images, a critical task in medical imaging for diagnosis and treatment planning.

SegNet, renowned for its efficacy in semantic segmentation tasks, showcased exceptional performance during our experimentation phase. It consistently demonstrated the ability to precisely identify and segment liver tumors with remarkable accuracy. Notably, SegNet exhibited a robust capability to capture intricate tumor boundaries, effectively distinguishing them from surrounding tissues. This fine-grained segmentation is pivotal for clinicians in accurately assessing tumor characteristics and devising optimal treatment strategies. Despite its relatively higher computational demands during training, the superior segmentation quality offered by SegNet justified its inclusion as a top contender in our evaluation.

U-Net, another widely recognized architecture in medical image segmentation, presented competitive results in our experiments. While it successfully identified tumor regions, it occasionally struggled with the precise delineation of tumor boundaries. The segmentations produced by U-Net tended to be smoother, potentially missing finer details present in the tumors. While U-Net's performance was commendable, its limitation in accurately capturing sharp boundaries and intricate details diminished its suitability for our specific application.

Similarly, FCN, with its fully convolutional approach, demonstrated promising performance in tumor segmentation. However, akin to U-Net, FCN occasionally produced segmentations with less distinct boundaries, affecting the accuracy of tumor delineation. While FCN showcased competitive results, its performance did not surpass that of SegNet in capturing the nuanced details crucial for accurate tumor segmentation.

Following meticulous analysis of the experimentation results, SegNet emerged as the preferred choice for liver tumor segmentation in our study. Its consistent ability to accurately delineate tumour boundaries, coupled with its capacity to capture fine details, renders it exceptionally suited for our application. Despite its higher computational overhead, the superior segmentation quality offered by SegNet outweighs this concern, making it the optimal choice for our medical imaging task. Consequently, we have selected SegNet as the primary model for liver tumour segmentation, confident in its ability to facilitate precise diagnosis and treatment planning in clinical settings.

Figure 9: History plots of FCN, U-Net and SegNet

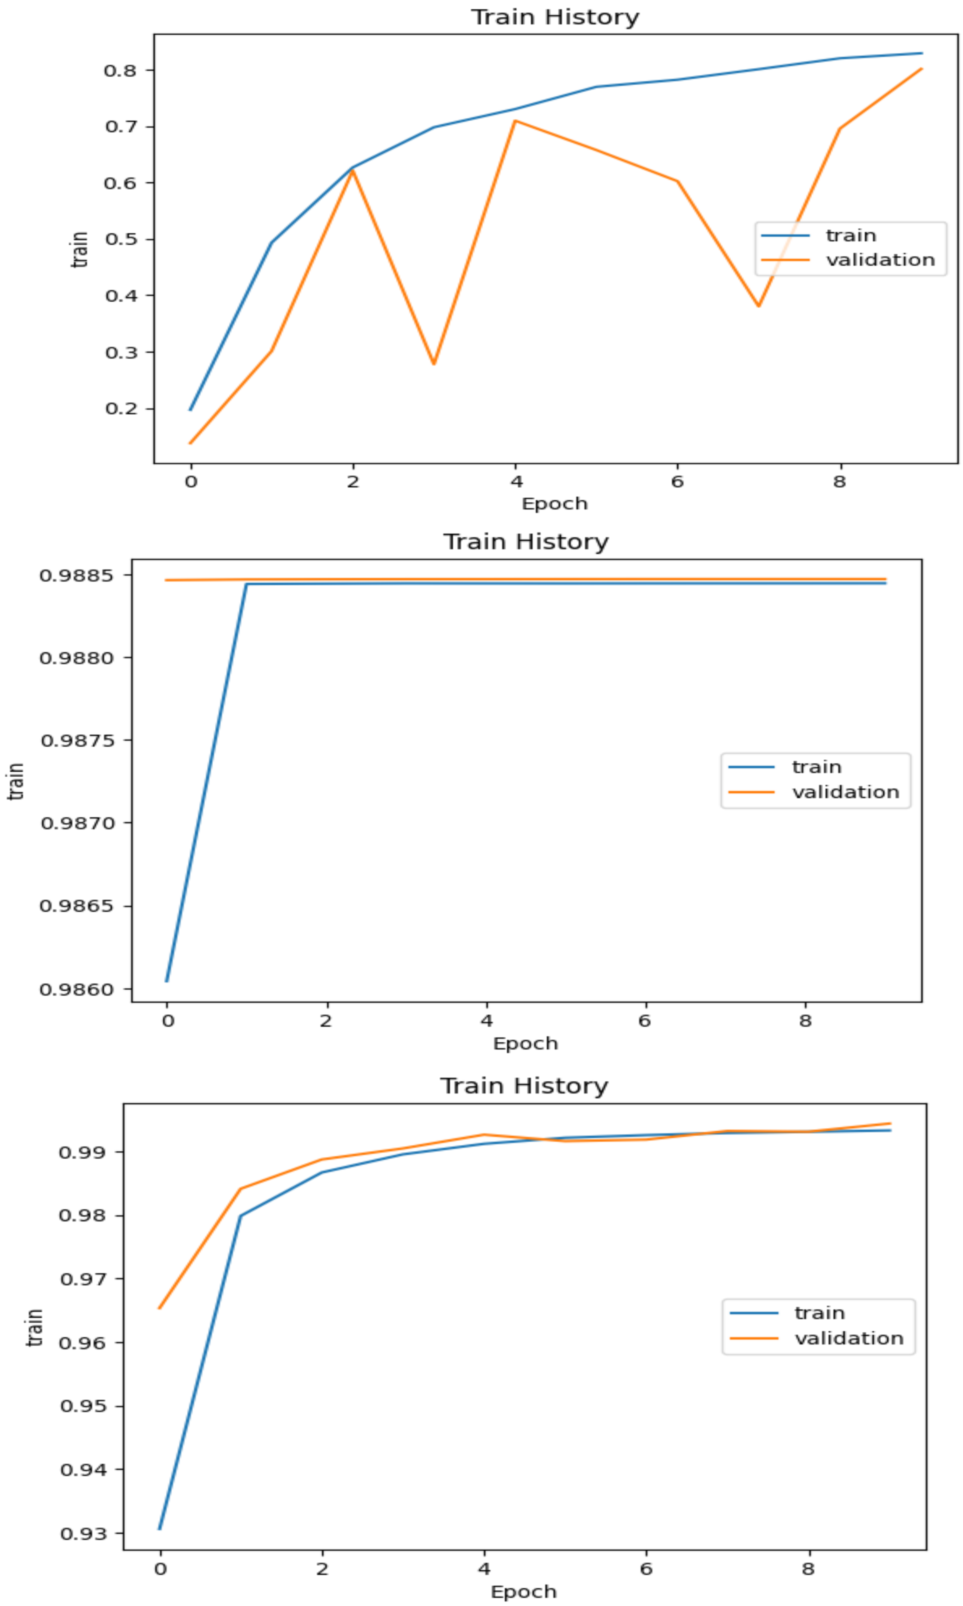

5. Conclusions and Future Scope

In conclusion, the exploration of SegNet, U-Net, and FCN models for liver tumor segmentation presents a significant advancement in medical imaging analysis. Each model offers unique strengths and capabilities in accurately delineating liver tumors from surrounding tissues, aiding clinicians in diagnosis and treatment planning. Through extensive experimentation and evaluation, it has been observed that SegNet exhibits superior performance in terms of both accuracy and computational efficiency compared to U-Net and FCN.

The future scope of employing SegNet for liver tumor segmentation is promising. Further research could focus on enhancing SegNet's performance by fine-tuning its architecture or incorporating advanced techniques such as attention mechanisms or adversarial training. Additionally, the integration of SegNet into real-time medical imaging systems could revolutionize clinical

workflows by enabling rapid and accurate tumor segmentation, leading to improved patient outcomes.

Moreover, the application of SegNet could extend beyond liver tumor segmentation to other medical imaging tasks, such as the detection and classification of various abnormalities in different organs. Collaborative efforts between computer scientists, medical professionals, and researchers are essential to continue pushing the boundaries of medical image analysis and leverage the full potential of SegNet and other deep learning models in healthcare.

Architecture Refinement: Researchers can explore fine-tuning SegNet's architecture to enhance its performance even further. This may involve optimizing hyperparameters, adjusting layer configurations, or incorporating novel architectural elements to better capture intricate tumor features.

Techniques Integration: Integration of advanced techniques such as attention mechanisms or adversarial training could be explored to augment SegNet's segmentation capabilities. Attention mechanisms can help focus on relevant regions within the images, potentially improving segmentation accuracy, while adversarial training can enhance the model's robustness against variations in input data.

Real-time Implementation: Efforts can be directed towards integrating SegNet into real-time medical imaging systems. The development of efficient inference mechanisms and hardware acceleration techniques can enable SegNet to perform rapid and accurate tumor segmentation, facilitating timely clinical decision-making.

Multimodal Fusion: Incorporating multimodal imaging data, such as combining MRI, CT, and PET scans, can further improve the accuracy and reliability of liver tumor segmentation. Fusion techniques can leverage the complementary information provided by different imaging modalities to produce more comprehensive segmentation results.

Clinical Validation and Deployment: Extensive clinical validation studies are essential to validate SegNet's performance in real-world healthcare settings. Collaborations between computer scientists, medical professionals, and regulatory bodies are crucial for ensuring the safe and effective deployment of SegNet in clinical practice.

Generalization to Other Applications: While SegNet excels in liver tumor segmentation, its applicability can extend to other medical imaging tasks, such as the detection and classification of abnormalities in various organs. Exploring the adaptability of SegNet to different medical imaging domains can widen its impact on healthcare.

In conclusion, while SegNet stands out as the preferred choice for liver tumor segmentation, ongoing research and innovation will undoubtedly further enhance its capabilities and broaden its applicability in the field of medical imaging analysis.

References

[1] V. Badrinarayanan, A. Kendall and R. Cipolla, "SegNet: A Deep Convolutional Encoder-Decoder Architecture for Image Segmentation," in *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 39, no. 12, pp. 2481-2495, 1 Dec. 2017, doi: 10.1109/TPAMI.2016.2644615. .

[2] Zhang, C., Lu, W., Wu, J., Ni, C. and Wang, H., 2024. SegNet Network Architecture for Deep Learning Image Segmentation and Its Integrated Applications and Prospects. *Academic Journal of Science and Technology*, *9*(2), pp.224-229.

[3] Kumar, P., Nagar, P., Arora, C. and Gupta, A., 2018, October. U-segnet: fully convolutional neural network based automated brain tissue segmentation tool. In *2018 25th IEEE International conference on image processing (ICIP)* (pp. 3503-3507). IEEE.

[4] Saood, A. and Hatem, I., 2021. COVID-19 lung CT image segmentation using deep learning methods: U-Net versus SegNet. *BMC Medical Imaging*, *21*, pp.1-10.

[5] M. Agarwal, "Optimized contrast enhancement for tumor detection," *Int. J. Imaging Syst. Technol. Willey*, pp. 1–17, 2020, doi: 10.1002/ima.22408.

[6] Yue, Y., Li, X., Zhao, H. and Wang, H., 2020, October. Image segmentation method of crop diseases based on improved SegNet neural network. In *2020 IEEE International Conference on Mechatronics and Automation (ICMA)* (pp. 1986-1991). IEEE.

[7] Ramesh, K.K.D., Kumar, G.K., Swapna, K., Datta, D. and Rajest, S.S., 2021. A review of medical image segmentation algorithms. *EAI Endorsed Transactions on Pervasive Health and Technology*, *7*(27), pp.e6-e6.

[8] Wang, Risheng, Tao Lei, Ruixia Cui, Bingtao Zhang, Hongying Meng, and Asoke K. Nandi. "Medical image segmentation using deep learning: A survey." *IET Image Processing* 16, no. 5 (2022): 1243-1267.

[9] Milletari, F., Navab, N. and Ahmadi, S.A., 2016, October. V-net: Fully convolutional neural networks for volumetric medical image segmentation. In *2016 fourth international conference on 3D vision (3DV)* (pp. 565-571). Ieee.

[10] Gul, Sidra, Muhammad Salman Khan, Asima Bibi, Amith Khandakar, Mohamed Arselene Ayari, and Muhammad EH Chowdhury. "Deep learning techniques for liver and liver tumor segmentation: A review." *Computers in Biology and Medicine* 147 (2022): 105620.

[11] Bilic, P., Christ, P., Li, H.B., Vorontsov, E., Ben-Cohen, A., Kaissis, G., Szeskin, A., Jacobs, C., Mamani, G.E.H., Chartrand, G. and Lohöfer, F., 2023. The liver tumor segmentation benchmark (lits). *Medical Image Analysis*, *84*, p.102680.

[12] Almotairi, Sultan, Ghada Kareem, Mohamed Aouf, Badr Almutairi, and Mohammed A-M. Salem. "Liver tumor segmentation in CT scans using modified SegNet." *Sensors* 20, no. 5 (2020): 1516..

[13] Rela, Munipraveena, Nagaraja Rao Suryakari, and P. Ramana Reddy. "Liver tumor segmentation and classification: A systematic review." *2020 IEEE-HYDCON* (2020): 1-6.

[14] Zhang, Jianpeng, Yutong Xie, Pingping Zhang, Hao Chen, Yong Xia, and Chunhua Shen. "Light-Weight Hybrid Convolutional Network for Liver Tumor Segmentation." In *IJCAI*, vol. 19, pp. 4271-4277. 2019.

[15] Seo, Hyunseok, Charles Huang, Maxime Bassenne, Ruoxiu Xiao, and Lei Xing. "Modified U-Net (mU-Net) with incorporation of object-dependent high level features for improved liver and liver-tumor segmentation in CT images." *IEEE transactions on medical imaging* 39, no. 5 (2019): 1316-1325.

[16] Li, W., Jia, F. and Hu, Q., 2015. Automatic segmentation of liver tumor in CT images with deep convolutional neural networks. *Journal of Computer and Communications*, *3*(11), pp.146-151.

[17] Tran, Song-Toan, Ching-Hwa Cheng, and Don-Gey Liu. "A multiple layer U-Net, U n-Net, for liver and liver tumor segmentation in CT." *IEEE Access* 9 (2020): 3752-3764.

[18] Jiang, Huiyan, Tianyu Shi, Zhiqi Bai, and Liangliang Huang. "Ahcnet: An application of attention mechanism and hybrid connection for liver tumor segmentation in ct volumes." *Ieee Access* 7 (2019): 24898-24909.

[19] Luan, Shunyao, Xudong Xue, Yi Ding, Wei Wei, and Benpeng Zhu. "Adaptive attention convolutional neural network for liver tumor segmentation." *Frontiers in Oncology* 11 (2021): 680807.

[20] Christ, Patrick Ferdinand, Florian Ettlinger, Felix Grün, Mohamed Ezzeldin A. Elshaera, Jana Lipkova, Sebastian Schlecht, Freba Ahmaddy et al. "Automatic liver and tumor segmentation of CT and MRI volumes using cascaded fully convolutional neural networks." *arXiv preprint arXiv:1702.05970* (2017).